Homework 2

Aiden Kenny

STAT GR5204: Statistical Inference

Columbia University

Novermber 26, 2020

Question 1

Question 9

If $\boldsymbol{X} \overset{\text{iid}}{\sim} \mathrm{N}(\mu, \sigma^2)$ with unknown μ and σ^2, then a $\gamma\%$ confidence interval for μ is given by

$$\mathcal{I} = \Big(\ \bar{X} - t_\gamma(n) \cdot S/\sqrt{n} \ , \ \bar{X} + t_\gamma(n) \cdot S/\sqrt{n} \ \Big),$$

where $t_\gamma(n) = T_{n-1}^{-1}((1+\gamma)/2)$ is the $(1+\gamma)/2$th quantile of the t distribution with $\mathrm{df} = n - 1$ and S is the sample standard deviation. The length of this confidence interval is given by

$$\Delta = \max(\mathcal{I}) - \min(\mathcal{I}) = \Big(\bar{X} + t_\gamma(n) \cdot S/\sqrt{n} \Big) - \Big(\bar{X} - t_\gamma(n) \cdot S/\sqrt{n} \Big) = 2t_\gamma(n) \cdot S/\sqrt{n}.$$

The squared length is then given by $\Delta^2 = 4t_\gamma^2(n) \cdot S^2/n$. Because the sample variance is an unbiased estimator for σ^2, we have $\mathbb{E}[\Delta^2] = \mathbb{E}\big[4t_\gamma^2(n) \cdot S^2/n\big] = 4t_\gamma^2(n) \cdot \sigma^2/n$. We now set $\mathbb{E}[\Delta^2] < \sigma^2/2$, and after some cancellations, we see that we need $t_\gamma^2(n)/n < 1/8$. There is no way to find a closed-form expression for this, so we will have to check the value of $t_\gamma^2(n)/n$ for increasing values of n. I set up a `while` loop in `R` to solve for it, and when $\gamma = 0.9$, we find that $n = 24$ is the smallest value of n such that $\mathbb{E}[\Delta^2] < \sigma^2/2$.

Question 10

Let $\boldsymbol{X} \overset{\text{iid}}{\sim} \mathrm{N}(\theta, \sigma^2)$, where θ is unknown and σ^2 is known, and we assume prior that $\theta \sim \mathrm{N}(\mu, \nu^2)$, where both μ and ν^2 are known.

(a) Since normal distributions are are conjugate to normal sampling, it follows that $\theta \,|\, \mathbf{x} \sim \mathrm{N}(\tilde{\mu}, \tilde{\sigma}^2)$, where

$$\tilde{\mu} = \frac{\sigma^2 \mu + n\nu^2 \bar{x}}{\sigma^2 n\nu^2} \quad \text{and} \quad \tilde{\sigma}^2 = \frac{\sigma^2 \nu^2}{\sigma^2 + n\mu^2}.$$

We also know that $(\theta \,|\, \mathbf{x} - \tilde{\mu})/\tilde{\sigma} \sim \mathrm{N}(0, 1)$, and so a 95% confidence interval for $\theta \,|\, \mathbf{x}$ is given by

$$\mathcal{I} = \Big(\ \tilde{\mu} - \Phi^{-1}(0.975) \cdot \tilde{\sigma} \ , \ \tilde{\mu} + \Phi^{-1}(0.975) \cdot \tilde{\sigma} \ \Big).$$

(b) We can think of our interval $\mathcal{I}$ as a function of ν^2. To examine what happens to $\mathcal{I}(\nu^2)$ as $\nu^2 \to \infty$, we will first look at $\tilde{\mu}$ and $\tilde{\sigma}$. Using L'Hopital's rule, we have

$$\lim_{\nu^2 \to \infty} \tilde{\mu} = \lim_{\nu^2 \to \infty} \frac{\sigma^2 \mu + n\nu^2 \bar{x}}{\sigma^2 n\nu^2} = \lim_{\nu^2 \to \infty} \frac{n\bar{x}}{n} = \bar{x},$$

$$\lim_{\nu^2 \to \infty} \tilde{\sigma} = \lim_{\nu^2 \to \infty} \sqrt{\frac{\sigma^2 \nu^2}{\sigma^2 + n\mu^2}} = \sqrt{\lim_{\nu^2 \to \infty} \frac{\sigma^2 \nu^2}{\sigma^2 + n\mu^2}} = \sqrt{\lim_{\nu^2 \to \infty} \frac{\sigma^2}{n}} = \frac{\sigma}{\sqrt{n}},$$

and so $\mathcal{I}(\nu^2) \to \big(\ \bar{x} - \Phi^{-1}(0.975) \cdot \sigma/\sqrt{n} \ , \ \bar{x} + \Phi^{-1}(0.975) \cdot \sigma/\sqrt{n} \ \big)$, which is a 95% confidence interval for θ.

Question 11

Let $\boldsymbol{X} \overset{\text{iid}}{\sim} \text{Unif}(0, \theta)$, where θ is unknown, and let $Y = X_{(n)}$ be the nth order statistic.

(a) Let $F_i(x) = \Pr(X_i \leq x) = x/\theta$ for all $i \in \{1, \ldots, n\}$. Then the cdf of Y is $G(y) = \prod_{i=1}^{n} F_i(y) = (y/\theta)^n$, since all of the X_i's are independent. The density of Y is then given by $g(y) = ny^{n-1}/\theta^n$ for $y \in [0, \theta]$. By letting $W = Y/\theta$, we have $Y = \theta W$ and $\partial Y/\partial W = \theta$, so the density of Y/θ is given by $h(w) = g\big(y(w)\big) \cdot \theta = nw^{n-1}$ for $w \in [0, 1]$. The cdf is then seen to be $H(w) = w^n$, and so the quantile is given by $H^{-1}(w) = \sqrt[n]{w}$.

(b) Since we must have $y \leq \theta$, it is natural that Y will underestimate θ, and is therefore biased. We have

$$\mathbb{E}[Y] = \int_0^\theta y \cdot \frac{ny^{n-1}}{\theta^n}\, \partial y = \frac{n}{\theta^n} \int_0^\theta y^n\, \partial y = \frac{n}{\theta^n} \cdot \frac{y^{n+1}}{n+1}\bigg|_0^\theta = \frac{n}{\theta^n} \cdot \frac{\theta^{n+1}}{n+1} = \frac{n\theta}{n+1},$$

and so the bias is $\text{bias}(Y) = \mathbb{E}[Y] - \theta = -\theta/(n+1)$.

(d) Using the cdf of Y/θ, for any interval involving Y/θ, we have $\Pr(a \leq Y/\theta \leq b) = b^n - a^n$, where $a, b \in (0, 1]$. Rearranging the terms inside the interval gives us $\Pr(Y/b \leq \theta \leq Y/a) = b^2 - a^2 \overset{\text{set}}{=} \gamma$. That is, as long as we impose the constraint that $b^2 - a^2 = \gamma$, any interval $(Y/b, Y/a)$ is a $\gamma\%$ confidence interval for θ.

Question 12

Suppose that $X \sim \text{P}$, where P is an unknown distribution, and we want to test $H_0 : \text{P} = \text{Unif}(0, 1)$ against $H_1 : \text{P} = \text{N}(0, 1)$. The densities of both distributions (within their support) are given by $f_0(x) = 1$ for $x \in [0, 1]$ and $f_2(x) = (2\pi)^{-1/2}\exp(-x^2/2)$, respectively, and so our test statistics is