Homework 3

Aiden Kenny
STAT GR5204: Statistical Inference
Columbia University

December 14, 2020

Question 1

Suppose that $\boldsymbol{X} \overset{\text{iid}}{\sim} \chi^2(\theta)$, where $\theta \in \mathbb{N}$ is unknown. We would like to test $H_0 : \theta \leq 8$ against $H_A : \theta > 8$, using a UMP test δ^* with a specified significance $\alpha_* \in (0, 1)$. The joint density of $\boldsymbol{X}$ is given by

$$f(\mathbf{x} \,|\, \theta) = \prod_{i=1}^{n} \frac{x_i^{\theta/2-1} \mathrm{e}^{-x/2}}{2^{\theta/2}\Gamma(\theta/2)} = 2^{-n\theta/2} \cdot \Gamma^{-n}(\theta/2) \cdot \left(\prod_{i=1}^{n} x_i\right)^{n(\theta/2-1)} \cdot \exp\left(-\frac{1}{2}\sum_{i=1}^{n} x_i\right)$$

To determine δ^*, we will look at the likelihood ratio. If we have two values θ_1, θ_2 such that $\theta_1 < \theta_2$, then then likelihood ratio is

$$\begin{aligned}\frac{f(\mathbf{x} \,|\, \theta_2)}{f(\mathbf{x} \,|\, \theta_1)} &= \frac{2^{-n\theta_1/2} \cdot \Gamma^{-n}(\theta_2/2) \cdot \left(\prod_i x_i\right)^{n(\theta_2/2-1)} \cdot \exp\left(-\frac{1}{2}\sum_i x_i\right)}{2^{-n\theta_1/2} \cdot \Gamma^{-n}(\theta_1/2) \cdot \left(\prod_i x_i\right)^{n(\theta_1/2-1)} \cdot \exp\left(-\frac{1}{2}\sum_i x_i\right)} \\ &= 2^{n(\theta_1-\theta_2)/2} \left(\frac{\Gamma(\theta_1/2)}{\Gamma(\theta_2/2)}\right)^n \left(\prod_{i=1}^{n} x_i\right)^{n(\theta_2-\theta_1)/2},\end{aligned}$$

which is a monotone increasing function of the test statistic $\prod_{i=1}^{n} x_i$. Therefore, the UMP test is δ^* : reject H_0 if $\prod_{i=1}^{n} X_i \geq c_*$, where c_* is chosen such that the test has a significance level α_*. Taking the log of both sides gives us $\sum_{i=1}^{8} \log X_i \geq \log c_* \coloneqq k$.

Question 2

Question 3

From a random sample of n people, we observe which of the k different cereal brands that each person prefers. Let N_i be the number of people who prefer the ith cereal (so $\sum_i N_i = n$), and let p_i be the proportion of people who prefer cereal i. With a significance of α, we would like to test $H_0 : p_1 = \cdots = p_k = 1/k$ against $H_A : H_0$ is false. To conduct this test we will use the χ^2 goodness-of-fit test, from which our test statistic is

$$Q = \sum_{i=1}^{k} \frac{(N_i - n/k)^2}{n/k} = \frac{k}{n}\sum_{i=1}^{k}\left(N_i^2 - \frac{2nN_i}{k} + \frac{n^2}{k^2}\right) = \frac{k}{n}\sum_{i=1}^{k} N_i^2 - n.$$

For this test, we reject H_0 when $Q \geq q_\alpha$, where q_α is the $1-\alpha$ quantile of the χ^2 distribution with $k-1$ degrees of freedom. As a result, we will reject H_0 when $\sum_i N_i^2 \geq n(q_\alpha + n)/k$. From the specific example we have $n = 400$, $k = 5$, and $q_{0.01} = 13.273$, so in order to reject H_0, we would need $\sum_i N_i^2 \geq 33062.136$.

Question 4

With significance α, we are testing whether or not each of the r groups have the same proportion for each of the c bloodtypes, that is, we are testing $H_0 : p_{1_j} = \cdots, = p_{r_j}$ for all $j = \{1, \ldots, c\}$ against $H_A : H_0$ is false (in this example, we have $2 \leq c \leq 4$). There are n total observations in the data broken down as follows: N_{ij} is the observation in the ith group and jth blood type, N_{i*} is the number of observations in the ith group across all c bloodtypes, and N_{*j} is the number of observations in the jth bloodtype across all r groups. Using the χ^2 test for homogenity, our test statistic is

$$Q = \sum_{i=1}^{r} \sum_{j=1}^{c} \frac{(N_{ij} - \hat{E}_{ij})^2}{\hat{E}_{ij}}, \quad \text{where} \quad \hat{E}_{ij} = \frac{N_{i*} \cdot N_{*j}}{n}.$$

We would the reject H_0 when $Q \geq q_\alpha$, where q_α is the $1 - \alpha$ quantile of the χ^2 distribution with $(r-1)(c-1)$ degrees of freedom. For this example, the contingency table is given to us, where we have $r = 3$, $c = 4$, and we want to use a significance level of $\alpha = 0.1$. Calculating the rest of the necessary values to obtain Q by hand is tediuos and will not be written here. Using `R`, I found that $Q = 6.829$ and $q_{0.1} = 10.645$, so we would not reject H_0. That is, the proportion of the four blood types seem to be the same for each of the three groups.

Question 5

Question 6

For this question, for all $i, j = \{1, 2\}$ we have $N_{i*} = N_{*j} = 2n$ and $\hat{E}_{ij} = 2n \cdot 2n/4n = n$, and so the test statistic for the χ^2 test of independence is

$$Q = \sum_{i=1}^{2} \sum_{j=1}^{2} \frac{(N_{ij} - \hat{E}_{ij})^2}{\hat{E}_{ij}} = \frac{4a^2}{n}.$$

We reject H_0 if $Q \geq q_\alpha$, where q_α is the $1 - \alpha$ quantile of the χ^2 distribution with 1 degree of freedom, and so we would reject H_0 when $a \geq \sqrt{nq_\alpha}/2$ or $a \leq -\sqrt{nq_\alpha}/2$. When $\alpha = 0.01$, we have $q_{0.01} = 6.635$.

Question 8

By definition, the least-squares estimate for the intercept is $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$. Rearranging gives us $\bar{y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{x}$, and so the least-squares line $y = \hat{\beta}_0 + \hat{\beta}_1 x$ will always pass through the point $(\bar{x}, \bar{y})$.

Question 9

(a) The least-squares coefficients for the model are given by $\hat{\beta}_0 = 40.893$ and $\hat{\beta}_1 = 0.548$.

Question 10

Let $(\mathbf{x}, \mathbf{y})$ be the vectors of observations for the predictor x and the response Y. From the data, we have $n = 10$, $\bar{x} = 2.33$, $\bar{y} = 0.81$, $\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 = 36.081$, and $(\mathbf{y} - \bar{y}\mathbf{1})^T(\mathbf{x} - \bar{x}\mathbf{1}) = 24.717$. Here we are assuming that $Y = \beta_0 + \beta_1 x + \epsilon$, where $\epsilon \sim \mathrm{N}(0, \sigma^2)$.

(a) The MLEs for β_0, β_1, and σ^2 are given by

$$\hat{\beta}_1 = \frac{(\mathbf{y} - \bar{y}\mathbf{1})^T(\mathbf{x} - \bar{x}\mathbf{1})}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} = 0.685, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = -0.786, \quad \hat{\sigma}^2 = \frac{\|\mathbf{y} - \hat{\beta}_0\mathbf{1} - \hat{\beta}_1\mathbf{x}\|^2}{n} = 0.938.$$

(b) The variance of $\hat{\beta}_0$ and $\hat{\beta}_1$ is given by

$$\mathrm{Var}\big[\hat{\beta}_1\big] = \frac{\sigma^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} = 0.0277\sigma^2, \quad \mathrm{Var}\big[\hat{\beta}_0\big] = \sigma^2\left(\frac{1}{n} + \frac{\bar{x}^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2}\right) = 0.25\sigma^2.$$

(c) The covariance between $\hat{\beta}_0$ and $\hat{\beta}_1$, and therefore the correltaion, is

$$\mathrm{Cov}\big[\hat{\beta}_0, \hat{\beta}_1\big] = -\frac{\bar{x}\sigma^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} = -0.0646\sigma^2, \quad \mathrm{Cor}\big[\hat{\beta}_0, \hat{\beta}_1\big] = \frac{\mathrm{Cov}\big[\hat{\beta}_0, \hat{\beta}_1\big]}{\sqrt{\mathrm{Var}\big[\hat{\beta}_0\big] \cdot \mathrm{Var}\big[\hat{\beta}_1\big]}} = -0.775.$$

Question 11

Suppose β_0, β_1 are the coefficients from the linear model in question 10, and we want to estimate $\theta = 3\beta_0 - 2\beta_1 + 5$. Because $\hat{\beta}_0$ and $\hat{\beta}_1$ are unbiased estimators for the coefficients, we can estimate θ with $\hat{\theta} = 3\hat{\beta}_0 - 2\hat{\beta}_1 + 5$. The MSE of $\hat{\theta}$, which is just its variance, is given by

$$\mathbb{E}\Big[(\hat{\theta} - \theta)^2\Big] = \mathrm{Var}\big[3\hat{\beta}_0 - 2\hat{\beta}_1 + 5\big] = 9\mathrm{Var}\big[\hat{\beta}_0\big] + 4\mathrm{Var}\big[\hat{\beta}_1\big] - 12\mathrm{Cov}\big[\hat{\beta}_0, \hat{\beta}_1\big] = 2.0245\sigma^2.$$

Question 12

We know that the MLE and least-squares estimates of β_0 and β_1 are the same, so $\hat{\beta}_0 = -0.786$ and $\hat{\beta}_1 = 0.685$. Using this linear model, when $x = 2$, we predict $\hat{Y} = -0.786_0.685 \cdot 2 = 0.584$. The variance of $\hat{Y}$ is given by

$$\mathbb{E}\Big[(\hat{Y} - Y)^2\Big] = \sigma^2\left(1 + \frac{1}{n} + \frac{(2 - \bar{x})^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2}\right) = 1.103\sigma^2.$$

Question 13